

Data Analysis

We will work with model Diabetes130US from openml with id 45069, first we will load this dataset.

```
import openml
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import torch

dataset_id = 45069

dataset = openml.datasets.get_dataset(dataset_id)
X, y, _, _ = dataset.get_data(dataset_format="dataframe",
target=dataset.default_target_attribute)

if y is not None:
    df = pd.concat([X, y], axis=1)
else:
    df = X

df.head()

   admission_type_id  discharge_disposition_id  admission_source_id \
0                  1                           3                         7
1                  5                           1                         1
2                  2                           1                         1
3                  3                           1                         1
4                  3                           1                         1

   time_in_hospital  num_lab_procedures  num_procedures
num_medications \
0                   4                      41                      2
17
1                   3                      64                      0
10
2                   2                      39                      0
8
3                   1                      39                      2
12
4                   3                      21                      2
23

   number_outpatient  number_emergency  number_inpatient  ...

```

```

citoglipiton \
0          0          0          0 ...
No
1          0          0          1 ...
No
2          0          0          0 ...
No
3          0          0          1 ...
No
4          1          0          0 ...
No

insulin glyburide.metformin glipizide.metformin
glimepiride.pioglitazone \
0 Steady          No          No
No
1 Steady          No          No
No
2 Steady          No          No
No
3 No             No          No
No
4 No             No          No
No

metformin.rosiglitazone metformin.pioglitazone change diabetesMed
class
0          No          No      Ch      Yes
No
1          No          No      Ch      Yes
No
2          No          No      Ch      Yes
No
3          No          No      No      No
No
4          No          No      Ch      Yes
No

[5 rows x 48 columns]

```

First we load the dataset from OpenML (if we run this for the first time we also download OpenML with pip).

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 48 columns):
 #   Column            Non-Null Count  Dtype 

```

0	admission_type_id	101766	non-null	uint8	
1	discharge_disposition_id	101766	non-null	uint8	
2	admission_source_id	101766	non-null	uint8	
3	time_in_hospital	101766	non-null	uint8	
4	num_lab_procedures	101766	non-null	uint8	
5	num_procedures	101766	non-null	uint8	
6	num_medications	101766	non-null	uint8	
7	number_outpatient	101766	non-null	uint8	
8	number_emergency	101766	non-null	uint8	
9	number_inpatient	101766	non-null	uint8	
10	number_diagnoses	101766	non-null	uint8	
11	race	99493	non-null	category	
12	gender	101766	non-null	category	
13	age	101766	non-null	category	
14	weight	3197	non-null	category	
15	payer_code	61510	non-null	category	
16	medical_specialty	51817	non-null	category	
17	diag_1	101745	non-null	category	
18	diag_2	101408	non-null	category	
19	diag_3	100343	non-null	category	
20	max_glu_serum	101766	non-null	category	
21	A1Cresult	101766	non-null	category	
22	metformin	101766	non-null	category	
23	repaglinide	101766	non-null	category	
24	nateglinide	101766	non-null	category	
25	chlorpropamide	101766	non-null	category	
26	glimepiride	101766	non-null	category	
27	acetohexamide	101766	non-null	category	
28	glipizide	101766	non-null	category	
29	glyburide	101766	non-null	category	
30	tolbutamide	101766	non-null	category	
31	pioglitazone	101766	non-null	category	
32	rosiglitazone	101766	non-null	category	
33	acarbose	101766	non-null	category	
34	miglitol	101766	non-null	category	
35	troglitazone	101766	non-null	category	
36	tolazamide	101766	non-null	category	
37	examide	101766	non-null	category	
38	citoglipiton	101766	non-null	category	
39	insulin	101766	non-null	category	
40	glyburide.metformin	101766	non-null	category	
41	glipizide.metformin	101766	non-null	category	
42	glimepiride.pioglitazone	101766	non-null	category	
43	metformin.rosiglitazone	101766	non-null	category	
44	metformin.pioglitazone	101766	non-null	category	
45	change	101766	non-null	category	
46	diabetesMed	101766	non-null	category	
47	class	101766	non-null	object	

```
dtypes: category(36), object(1), uint8(11)
memory usage: 5.7+ MB
```

We can see that we have 101 766 records with 48 features, 36 being categorical, 1 being object and 11 being integer.

```
df.unique()

admission_type_id           8
discharge_disposition_id    26
admission_source_id         17
time_in_hospital            14
num_lab_procedures          118
num_procedures                7
num_medications               75
number_outpatient              39
number_emergency                33
number_inpatient                 21
number_diagnoses                  16
race                           5
gender                          3
age                            10
weight                          9
payer_code                      17
medical_specialty                72
diag_1                          716
diag_2                          748
diag_3                          789
max_glu_serum                   4
A1Cresult                        4
metformin                         4
repaglinide                        4
nateglinide                        4
chlorpropamide                     4
glimepiride                        4
acetohexamide                      2
glipizide                           4
glyburide                           4
tolbutamide                          2
pioglitazone                        4
rosiglitazone                       4
acarbose                            4
miglitol                            4
troglitazone                         2
tolazamide                            3
examide                             1
citoglipton                           1
insulin                             4
glyburide.metformin                  4
```

```

glipizide.metformin      2
glimepiride.pioglitazone 2
metformin.rosiglitazone   2
metformin.pioglitazone    2
change                     2
diabetesMed                 2
class                      3
dtype: int64

```

We can see that examine and citoglipton have only 1 unique value, therefore we can drop them. We will also drop weight feature since it is filled in only 3197 times (which is about 3.14%).

```

columns_to_drop = ['examide', 'citoglipton', 'weight']

df = df.drop(columns=columns_to_drop)

print(f"Columns {columns_to_drop} have been dropped.")
print("\nDataFrame columns after dropping:")
print(df.columns)

Columns ['examide', 'citoglipton', 'weight'] have been dropped.

DataFrame columns after dropping:
Index(['admission_type_id', 'discharge_disposition_id',
       'admission_source_id',
       'time_in_hospital', 'num_lab_procedures', 'num_procedures',
       'num_medications', 'number_outpatient', 'number_emergency',
       'number_inpatient', 'number_diagnoses', 'race', 'gender',
       'age',
       'payer_code', 'medical_specialty', 'diag_1', 'diag_2',
       'diag_3',
       'max_glu_serum', 'A1Cresult', 'metformin', 'repaglinide',
       'nateglinide',
       'chlorpropamide', 'glimepiride', 'acetohexamide', 'glipizide',
       'glyburide', 'tolbutamide', 'pioglitazone', 'rosiglitazone',
       'acarbose',
       'miglitol', 'troglitazone', 'tolazamide', 'insulin',
       'glyburide.metformin', 'glipizide.metformin',
       'glimepiride.pioglitazone', 'metformin.rosiglitazone',
       'metformin.pioglitazone', 'change', 'diabetesMed', 'class'],
      dtype='object')

```

Next we will take features that represent ids and feature class and make it into categorical features.

```

columns_to_convert_to_category = ['admission_type_id',
       'discharge_disposition_id', 'admission_source_id', 'class']

for col in columns_to_convert_to_category:

```


```
df[col] = df[col].astype('category')

print("Converted columns to category:")
print(df[columns_to_convert_to_category].dtypes)

Converted columns to category:
admission_type_id      category
discharge_disposition_id      category
admission_source_id      category
class                      category
dtype: object

df.hist(figsize=(18, 18), bins=20)
plt.suptitle("Distributions of attributes")
plt.show()
```

Distributions of attributes



```
categorical_cols = df.select_dtypes(include='category').columns

print("Unique values for categorical features:")

for col in categorical_cols:
    print(f"\nColumn: {col}")
    print(df[col].unique())
```

```
Unique values for categorical features:
```

```
Column: admission_type_id
```

```
[1, 5, 2, 3, 6, 8, 4, 7]
```

```
Categories (8, uint8): [1, 2, 3, 4, 5, 6, 7, 8]
```

```
Column: discharge_disposition_id
```

```
[3, 1, 6, 22, 18, ..., 10, 17, 27, 19, 20]
```

```
Length: 26
```

```
Categories (26, uint8): [1, 2, 3, 4, ..., 24, 25, 27, 28]
```

```
Column: admission_source_id
```

```
[7, 1, 17, 6, 4, ..., 8, 22, 11, 25, 13]
```

```
Length: 17
```

```
Categories (17, uint8): [1, 2, 3, 4, ..., 17, 20, 22, 25]
```

```
Column: race
```

```
['AfricanAmerican', 'Caucasian', 'Hispanic', 'Other', NaN, 'Asian']
```

```
Categories (5, object): ['AfricanAmerican' < 'Asian' < 'Caucasian' < 'Hispanic' < 'Other']
```

```
Column: gender
```

```
['Female', 'Male', 'Unknown/Invalid']
```

```
Categories (3, object): ['Female' < 'Male' < 'Unknown/Invalid']
```

```
Column: age
```

```
['[60-70)', '[50-60)', '[80-90)', '[40-50)', '[70-80)', '[90-100)', '[10-20)', '[30-40)', '[20-30)', '[0-10)']
```

```
Categories (10, object): ['[0-10)' < '[10-20)' < '[20-30)' < '[30-40)' ... '[60-70)' < '[70-80)' < '[80-90)' < '[90-100)']
```

```
Column: payer_code
```

```
['MC', 'MD', NaN, 'BC', 'SP', ..., 'DM', 'WC', 'SI', 'OT', 'FR']
```

```
Length: 18
```

```
Categories (17, object): ['BC' < 'CH' < 'CM' < 'CP' ... 'SI' < 'SP' < 'UN' < 'WC']
```

```
Column: medical_specialty
```

```
[NaN, 'Surgery-Vascular', 'InternalMedicine', 'Surgery-General', 'Family/GeneralPractice', ..., 'Surgery-PlasticwithinHeadandNeck', 'Dermatology', 'Perinatology', 'Pediatrics-InfectiousDiseases', 'Psychiatry-Addictive']
```

```
Length: 73
```

```
Categories (72, object): ['AllergyandImmunology' < 'Anesthesiology' < 'Anesthesiology-Pediatric' < 'Cardiology' ... 'Surgery-Thoracic' < 'Surgery-Vascular' < 'SurgicalSpecialty' < 'Urology']
```

```
Column: diag_1
```

```
['820', '276', '493', '250.4', '218', ..., '133', '913', '827', '373', '837']
```

```
Length: 717
Categories (716, object): ['10' < '11' < '110' < '112' ... 'V66' <
'V67' < 'V70' < 'V71']

Column: diag_2
['599', '295', '250', '403', '627', ..., '915', '268', '665', '908',
'894']
Length: 749
Categories (748, object): ['11' < '110' < '111' < '112' ... 'V70' <
'V72' < 'V85' < 'V86']

Column: diag_3
['294', '780', NaN, '285', '620', ..., '195', '684', 'E864', '111',
'838']
Length: 790
Categories (789, object): ['11' < '110' < '111' < '112' ... 'V70' <
'V72' < 'V85' < 'V86']

Column: max_glu_serum
['None', 'Norm', '>300', '>200']
Categories (4, object): ['>200' < '>300' < 'None' < 'Norm']

Column: A1Cresult
['None', '>8', '>7', 'Norm']
Categories (4, object): ['>7' < '>8' < 'None' < 'Norm']

Column: metformin
['No', 'Steady', 'Down', 'Up']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: repaglinide
['No', 'Steady', 'Up', 'Down']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: nateglinide
['No', 'Steady', 'Down', 'Up']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: chlorpropamide
['No', 'Steady', 'Up', 'Down']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: glimepiride
['No', 'Steady', 'Down', 'Up']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: acetohexamide
['No', 'Steady']
Categories (2, object): ['No' < 'Steady']

Column: glipizide
```

```
['Steady', 'No', 'Up', 'Down']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: glyburide
['No', 'Steady', 'Up', 'Down']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: tolbutamide
['No', 'Steady']
Categories (2, object): ['No' < 'Steady']

Column: pioglitazone
['No', 'Steady', 'Up', 'Down']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: rosiglitazone
['No', 'Steady', 'Down', 'Up']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: acarbose
['No', 'Steady', 'Down', 'Up']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: miglitol
['No', 'Steady', 'Up', 'Down']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: troglitazone
['No', 'Steady']
Categories (2, object): ['No' < 'Steady']

Column: tolazamide
['No', 'Steady', 'Up']
Categories (3, object): ['No' < 'Steady' < 'Up']

Column: insulin
['Steady', 'No', 'Down', 'Up']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: glyburide.metformin
['No', 'Steady', 'Down', 'Up']
Categories (4, object): ['Down' < 'No' < 'Steady' < 'Up']

Column: glipizide.metformin
['No', 'Steady']
Categories (2, object): ['No' < 'Steady']

Column: glimepiride.pioglitazone
['No', 'Steady']
Categories (2, object): ['No' < 'Steady']
```

```
Column: metformin/rosiglitazone
['No', 'Steady']
Categories (2, object): ['No' < 'Steady']

Column: metformin/pioglitazone
['No', 'Steady']
Categories (2, object): ['No' < 'Steady']

Column: change
['Ch', 'No']
Categories (2, object): ['Ch' < 'No']

Column: diabetesMed
['Yes', 'No']
Categories (2, object): ['No' < 'Yes']

Column: class
['NO', '>30', '<30']
Categories (3, object): ['<30', '>30', 'NO']
```

Next we drop duplicates.

```
print("Shape before removing duplicates:", df.shape)

df.drop_duplicates(inplace=True)

print("Shape after removing duplicates:", df.shape)
```

Shape before removing duplicates: (101766, 45)
Shape after removing duplicates: (101766, 45)

We have no duplicates.

```
print(df.isnull().sum())
```

admission_type_id	0
discharge_disposition_id	0
admission_source_id	0
time_in_hospital	0
num_lab_procedures	0
num_procedures	0
num_medications	0
number_outpatient	0
number_emergency	0
number_inpatient	0
number_diagnoses	0
race	2273
gender	0
age	0

payer_code	40256
medical_specialty	49949
diag_1	21
diag_2	358
diag_3	1423
max_glu_serum	0
A1Cresult	0
metformin	0
repaglinide	0
nateglinide	0
chlorpropamide	0
glimepiride	0
acetohexamide	0
glipizide	0
glyburide	0
tolbutamide	0
pioglitazone	0
rosiglitazone	0
acarbose	0
miglitol	0
troglitazone	0
tolazamide	0
insulin	0
glyburide.metformin	0
glipizide.metformin	0
glimepiride.pioglitazone	0
metformin/rosiglitazone	0
metformin/pioglitazone	0
change	0
diabetesMed	0
class	0
dtype:	int64

We have null values in race, payer_code, medical_specialty, diag_1, diag_2, diag_3 columns, we will add missing value for race, payer_code and medical_specialty. For diagnoses we will make columns for each possible diagnosis value and mark it as 1 if its present in either of diag_1, diag_2 or diag_3, since it doesn't matter if the diagnosis is in diag_1 or diag_3.

```
categorical_cols_with_missing = ['race', 'payer_code',
'medical_specialty']

for col in categorical_cols_with_missing:
    if col in df.columns:
        if pd.api.types.is_categorical_dtype(df[col].dtype):
            if 'Missing' not in df[col].cat.categories:
                df[col] = df[col].cat.add_categories('Missing')
            df[col].fillna('Missing', inplace=True)
```

```
print("Missing values:")
display(df[categorical_cols_with_missing].isnull().sum())

Missing values:

C:\Users\maria\AppData\Local\Temp\ipykernel_26612\456492141.py:6:
DeprecationWarning: is_categorical_dtype is deprecated and will be
removed in a future version. Use isinstance(dtype,
pd.CategoricalDtype) instead
    if pd.api.types.is_categorical_dtype(df[col].dtype):
C:\Users\maria\AppData\Local\Temp\ipykernel_26612\456492141.py:9:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df[col].fillna('Missing', inplace=True)

race          0
payer_code    0
medical_specialty 0
dtype: int64

all_diag_codes = pd.concat([df['diag_1'], df['diag_2'],
                           df['diag_3']]).unique()

all_diag_codes = [code for code in all_diag_codes if pd.notna(code)]

print(f"Number of unique diagnosis: {len(all_diag_codes)}")

print(all_diag_codes[:10])

Number of unique diagnosis: 915
['820', '276', '493', '250.4', '218', '250.6', '410', '486', '584',
 '428']

diag_cols_dict = {
    f'diag|{code}': ((df['diag_1'] == code) |
                      (df['diag_2'] == code) |
                      (df['diag_3'] == code)).astype(int)
    for code in all_diag_codes
}
```

```

diag_cols_df = pd.DataFrame(diag_cols_dict, index=df.index)

df = df.drop(columns=['diag_1', 'diag_2', 'diag_3'])

df = pd.concat([df, diag_cols_df], axis=1)

print("New DataFrame:")
display(df.head())

```

New DataFrame:

	admission_type_id	discharge_disposition_id	admission_source_id	
0	1	3	7	
1	5	1	1	
2	2	1	1	
3	3	1	1	
4	3	1	1	

	time_in_hospital	num_lab_procedures	num_procedures	
0	4	41	2	
17				
1	3	64	0	
10				
2	2	39	0	
8				
3	1	39	2	
12				
4	3	21	2	
23				

	number_outpatient	number_emergency	number_inpatient	...	diag
744	0	0	0	0	...
0					
1	0	0	0	1	...
0					
2	0	0	0	0	...
0					
3	0	0	0	1	...
0					
4	1	0	0	0	...
0					

	diag V22	diag 265	diag E912	diag E922	diag E861	diag E865	diag
387	0	0	0	0	0	0	0
0							
1	0	0	0	0	0	0	0

2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0

	diag E966	diag E864
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 957 columns]

Next we will take a look at numerical features.

```
df.describe()

      time_in_hospital  num_lab_procedures  num_procedures
num_medications \
count    101766.000000          101766.000000  101766.000000
101766.000000
mean        4.395987           43.095641   1.339730
16.021844
std         2.985108           19.674362   1.705807
8.127566
min         1.000000           1.000000   0.000000
1.000000
25%         2.000000           31.000000   0.000000
10.000000
50%         4.000000           44.000000   1.000000
15.000000
75%         6.000000           57.000000   2.000000
20.000000
max        14.000000          132.000000   6.000000
81.000000

      number_outpatient  number_emergency  number_inpatient \
count    101766.000000          101766.000000  101766.000000
mean        0.369357           0.197836   0.635566
std         1.267265           0.930472   1.262863
min         0.000000           0.000000   0.000000
25%         0.000000           0.000000   0.000000
50%         0.000000           0.000000   0.000000
75%         0.000000           0.000000   1.000000
max        42.000000           76.000000  21.000000

      number_diagnoses  diag|820  diag|276  ...
diag|
```

```

744 \
count      101766.000000 101766.000000 101766.000000 ...
101766.000000
mean       7.422607      0.011212      0.127361 ...
0.000010
std        1.933600      0.105292      0.333378 ...
0.003135
min        1.000000      0.000000      0.000000 ...
0.000000
25%        6.000000      0.000000      0.000000 ...
0.000000
50%        8.000000      0.000000      0.000000 ...
0.000000
75%        9.000000      0.000000      0.000000 ...
0.000000
max        16.000000     1.000000      1.000000 ...
1.000000

          diag|V22      diag|265      diag|E912      diag|E922 \
count    101766.000000 101766.000000 101766.000000 101766.000000
mean      0.000010      0.000010      0.000010      0.000010
std       0.003135      0.003135      0.003135      0.003135
min      0.000000      0.000000      0.000000      0.000000
25%      0.000000      0.000000      0.000000      0.000000
50%      0.000000      0.000000      0.000000      0.000000
75%      0.000000      0.000000      0.000000      0.000000
max      1.000000      1.000000      1.000000      1.000000

          diag|E861      diag|E865      diag|387      diag|E966 \
count    101766.000000 101766.000000 101766.000000 101766.000000
mean      0.000010      0.000010      0.000010      0.000010
std       0.003135      0.003135      0.003135      0.003135
min      0.000000      0.000000      0.000000      0.000000
25%      0.000000      0.000000      0.000000      0.000000
50%      0.000000      0.000000      0.000000      0.000000
75%      0.000000      0.000000      0.000000      0.000000
max      1.000000      1.000000      1.000000      1.000000

          diag|E864
count    101766.000000
mean      0.000010
std       0.003135
min      0.000000
25%      0.000000
50%      0.000000
75%      0.000000
max      1.000000

[8 rows x 923 columns]

```

```


numerical_cols = ['admission_type_id', 'discharge_disposition_id',
'admission_source_id', 'time_in_hospital', 'num_lab_procedures',
'num_procedures', 'num_medications', 'number_outpatient',
'number_emergency', 'number_inpatient', 'number_diagnoses']
id_cols = ['admission_type_id', 'discharge_disposition_id',
'admission_source_id']

numerical_cols_for_plotting = [col for col in numerical_cols if col
not in id_cols]


print("Boxplots of selected numerical columns (excluding IDs):")
for col in numerical_cols_for_plotting:
    plt.figure(figsize=(8, 6))
    sns.boxplot(data=df, y=col)
    plt.title(f'Boxplot of {col}')
    plt.ylabel(col)
    plt.show()

```


Boxplots of selected numerical columns (excluding IDs):




Boxplot of num_lab_procedures




Boxplot of num_procedures




Boxplot of num_medications




Boxplot of number_outpatient




Boxplot of number_emergency



Boxplot of number_inpatient





```

print("Number of outliers (using IQR method):")

for col in numerical_cols_for_plotting:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]
    num_outliers = outliers.shape[0]

    print(f"Column '{col}': {num_outliers} outliers")

Number of outliers (using IQR method):
Column 'time_in_hospital': 2252 outliers
Column 'num_lab_procedures': 143 outliers
Column 'num_procedures': 4954 outliers
Column 'num_medications': 2557 outliers
Column 'number_outpatient': 16739 outliers

```

```
Column 'number_emergency': 11383 outliers
Column 'number_inpatient': 7049 outliers
Column 'number_diagnoses': 281 outliers
```

We can see there are many outliers, but we won't replace all of them. For number_inpatient, number_emergency, number_outpatient the distribution is very tight around one number, if we would replace the outliers in these features, all the values would be the same. We would lose potential information. Therefore we will deal with outliers in number_diagnoses, num_lab_procedures, num_medications, num_procedures and time_in_hospital only, we will find outliers using IQR and then we will set them to the value of quantile bound.

```
columns_to_cap_iqr = ['number_diagnoses', 'num_lab_procedures',
'num_medications', 'num_procedures', 'time_in_hospital']

for col in columns_to_cap_iqr:

    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    df[col] = df[col].apply(lambda x: lower_bound if x < lower_bound
else x)

    df[col] = df[col].apply(lambda x: upper_bound if x > upper_bound
else x)

print("\nDescriptive Statistics for capped columns after IQR
capping:")
display(df[columns_to_cap_iqr].describe())
```

Descriptive Statistics for capped columns after IQR capping:

	number_diagnoses	num_lab_procedures	num_medications
count	101766.000000	101766.000000	101766.000000
mean	7.422395	43.087210	15.808512
std	1.925325	19.648915	7.396645
min	1.500000	1.000000	1.000000
25%	6.000000	31.000000	10.000000
50%	8.000000	44.000000	15.000000

```
75%           9.000000      57.000000      20.000000  
2.000000  
max          13.500000      96.000000      35.000000  
5.000000
```

```
    time_in_hospital  
count     101766.000000  
mean       4.363618  
std        2.892181  
min        1.000000  
25%        2.000000  
50%        4.000000  
75%        6.000000  
max        12.000000
```

```
columns_to_plot_iqr_capped = ['number_diagnoses',  
'num_lab_procedures', 'num_medications', 'num_procedures',  
'time_in_hospital']
```


```
print("Boxplots for columns after IQR-based outlier capping:")
```

```
for col in columns_to_plot_iqr_capped:
```


```
    plt.figure(figsize=(8, 6))  
    sns.boxplot(data=df, y=col)  
    plt.title(f'Boxplot of {col} (IQR Capped)')  
    plt.ylabel(col)  
    plt.show()
```

```
Boxplots for columns after IQR-based outlier capping:
```


Boxplot of number_diagnoses (IQR Capped)




Boxplot of num_lab_procedures (IQR Capped)




Boxplot of num_medications (IQR Capped)



Boxplot of num_procedures (IQR Capped)





Now we can see there are no outliers in these selected features.

At last we will scale the features to be in range of -1,1.

```
from sklearn.preprocessing import MinMaxScaler

numerical_cols = ['time_in_hospital', 'num_lab_procedures',
'num_procedures', 'num_medications', 'number_outpatient',
'number_emergency', 'number_inpatient', 'number_diagnoses']

scaler = MinMaxScaler()

df[numerical_cols] = scaler.fit_transform(df[numerical_cols])

print("Scaled numerical features:")
display(df[numerical_cols].head())

Scaled numerical features:

    time_in_hospital  num_lab_procedures  num_procedures
num_medications \
0            0.272727           0.421053           0.4
0.470588
```

1	0.181818	0.663158	0.0
0.264706			
2	0.090909	0.400000	0.0
0.205882			
3	0.000000	0.400000	0.4
0.323529			
4	0.181818	0.210526	0.4
0.647059			
	number_outpatient	number_emergency	number_inpatient
	number_diagnoses		
0	0.00000	0.0	0.000000
0.625000			
1	0.00000	0.0	0.047619
0.458333			
2	0.00000	0.0	0.000000
0.041667			
3	0.00000	0.0	0.047619
0.125000			
4	0.02381	0.0	0.000000
0.625000			

Next step is to encode categories to be represented by numbers only.

```
categorical_cols_for_encoding =
df.select_dtypes(include=['category']).columns.tolist()

cols_to_keep_all = [col for col in categorical_cols_for_encoding if
df[col].nunique() > 2]

cols_to_drop_first = [col for col in categorical_cols_for_encoding if
df[col].nunique() == 2]

df = pd.get_dummies(df, columns=cols_to_drop_first, drop_first=True,
prefix_sep='|')

df = pd.get_dummies(df, columns=cols_to_keep_all, drop_first=False,
prefix_sep='|')


print("Shape of new DataFrame:", df.shape)
print("\nColumns after one-hot encoding:")
print(df.columns)

df.head()

Shape of new DataFrame: (101766, 1159)

Columns after one-hot encoding:
```

```
Index(['time_in_hospital', 'num_lab_procedures', 'num_procedures',
       'num_medications', 'number_outpatient', 'number_emergency',
       'number_inpatient', 'number_diagnoses', 'diag|820', 'diag|276',
       ...
       'insulin|No', 'insulin|Steady', 'insulin|Up',
       'glyburide.metformin|Down', 'glyburide.metformin|No',
       'glyburide.metformin|Steady', 'glyburide.metformin|Up', 'class|<30',
       'class|>30', 'class|No'],
      dtype='object', length=1159)
```

```
   time_in_hospital  num_lab_procedures  num_procedures
num_medications \
0           0.272727          0.421053          0.4
0.470588
1           0.181818          0.663158          0.0
0.264706
2           0.090909          0.400000          0.0
0.205882
3           0.000000          0.400000          0.4
0.323529
4           0.181818          0.210526          0.4
0.647059
```

```
   number_outpatient  number_emergency  number_inpatient
number_diagnoses \
0           0.00000          0.0          0.000000
0.625000
1           0.00000          0.0          0.047619
0.458333
2           0.00000          0.0          0.000000
0.041667
3           0.00000          0.0          0.047619
0.125000
4           0.02381          0.0          0.000000
0.625000
```

```
   diag|820  diag|276  ...  insulin|No  insulin|Steady  insulin|Up \
0       1      0  ...  False        True        False
1       0      1  ...  False        True        False
2       0      0  ...  False        True        False
3       0      0  ...  True        False        False
4       0      0  ...  True        False        False

   glyburide.metformin|Down  glyburide.metformin|No \
0           False          True
1           False          True
2           False          True
3           False          True
4           False          True
```

```

glyburide.metformin|Steady  glyburide.metformin|Up   class|<30
class|>30  \
0           False        False        False
False
1           False        False        False
False
2           False        False        False
False
3           False        False        False
False
4           False        False        False
False

class|NO
0    True
1    True
2    True
3    True
4    True

[5 rows x 1159 columns]

encoded_columns = df.columns

feature_counts = {}

for col in encoded_columns:
    parts = col.split('|')

    if len(parts) > 1:
        original_feature = parts[0]

        if original_feature in feature_counts:
            feature_counts[original_feature] += 1
        else:
            feature_counts[original_feature] = 1

print("Number of encoded columns per original feature:")
for feature, count in feature_counts.items():
    print(f"{feature}: {count}")

Number of encoded columns per original feature:
diag: 915
acetohexamide: 1
tolbutamide: 1
troglitazone: 1
glipizide.metformin: 1
glimepiride.pioglitazone: 1
metformin.rosiglitazone: 1

```

```
metformin.pioglitazone: 1
change: 1
diabetesMed: 1
admission_type_id: 8
discharge_disposition_id: 26
admission_source_id: 17
race: 6
gender: 3
age: 10
payer_code: 18
medical_specialty: 73
max_glu_serum: 4
A1Cresult: 4
metformin: 4
repaglinide: 4
nateglinide: 4
chlorpropamide: 4
glimepiride: 4
glipizide: 4
glyburide: 4
pioglitazone: 4
rosiglitazone: 4
acarbose: 4
miglitol: 4
tolazamide: 3
insulin: 4
glyburide.metformin: 4
class: 3
```

Now we have our dataset with numeric features only, but it has 1159 features. We will try and drop features that have either small correlation or have low importance.

```
correlation_matrix = df.corrwith(df['diabetesMed|Yes'])

encoded_cols = [col for col in df.columns if df[col].dtype != 'float']
df[encoded_cols] = df[encoded_cols].astype(int)

print("\nCorrelation with 'diabetesMed|Yes':")
display(correlation_matrix.sort_values(ascending=False).head(30))
display(correlation_matrix.sort_values(ascending=True).head(30))
```

Correlation with 'diabetesMed|Yes':

diabetesMed Yes	1.000000
insulin Steady	0.360434
metformin Steady	0.256281
insulin Down	0.201861
glipizide Steady	0.193680

num_medications	0.193321
insulin Up	0.193296
glyburide Steady	0.173046
pioglitazone Steady	0.148253
rosiglitazone Steady	0.137996
glimepiride Steady	0.119850
A1Cresult >8	0.093123
repaglinide Steady	0.064168
time_in_hospital	0.063169
diag 250.02	0.058306
metformin Up	0.056254
glyburide Up	0.049011
glipizide Up	0.047717
class >30	0.046396
glyburide.metformin Steady	0.045218
payer_code MC	0.044897
nateglinide Steady	0.044422
metformin Down	0.041195
glyburide Down	0.040797
glipizide Down	0.040651
medical_specialty Emergency/Trauma	0.038697
payer_code SP	0.034813
payer_code MD	0.033967
num_lab_procedures	0.033181
diag 250.13	0.031863
dtype: float64	
insulin No	-0.585464
change No	-0.506370
metformin No	-0.270176
glipizide No	-0.206230
glyburide No	-0.186835
pioglitazone No	-0.152230
rosiglitazone No	-0.141157
glimepiride No	-0.126699
payer_code Missing	-0.075663
discharge_disposition_id 18	-0.069842
A1Cresult None	-0.069683
repaglinide No	-0.067718
class NO	-0.061508
max_glu_serum Norm	-0.046610
glyburide.metformin No	-0.045677
nateglinide No	-0.045579
diag 560	-0.041975
discharge_disposition_id 11	-0.033247
diag 562	-0.032704
diag 789	-0.031176
admission_type_id 6	-0.030737
payer_code HM	-0.030696
acarbose No	-0.030110

```

age|[90-100)           -0.029543
admission_source_id|5   -0.027727
diag|578                -0.027317
diag|553                -0.024053
diag|574                -0.023991
medical_specialty|Missing -0.023135
diag|558                -0.022460
dtype: float64

import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split

X = df.drop(columns=['diabetesMed|Yes'])
y = df['diabetesMed|Yes']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

model = RandomForestClassifier(n_estimators=200, random_state=42,
n_jobs=-1)
model.fit(X_train, y_train)

importances = pd.Series(model.feature_importances_, index=X.columns)

importances_sorted = importances.sort_values(ascending=False)
print(importances_sorted.head(50))

insulin|No              0.213665
insulin|Steady          0.113936
change|No                0.099622
metformin|No             0.054329
metformin|Steady         0.046696
glipizide|No             0.037536
glipizide|Steady         0.036931
glyburide|No             0.035107
glyburide|Steady         0.031467
insulin|Down             0.024147
insulin|Up               0.021271
pioglitazone|No          0.017476
pioglitazone|Steady      0.017219
num_medications          0.016498
glimepiride|Steady       0.014636
glimepiride|No            0.014391
rosiglitazone|No          0.013940
rosiglitazone|Steady      0.013675
num_lab_procedures        0.008100
time_in_hospital          0.005927
number_diagnoses          0.004706
repaglinide|No            0.004435

```

```

num_procedures          0.003829
repaglinide|Steady      0.003725
number_inpatient         0.003229
A1Cresult|>8           0.002546
payer_code|Missing       0.002486
glyburide.metformin|No   0.002328
number_outpatient        0.002222
glyburide.metformin|Steady 0.002204
class|NO                 0.002180
nateglinide|No           0.001763
medical_specialty|Missing 0.001742
nateglinide|Steady       0.001690
diag|250                 0.001687
number_emergency         0.001676
age|[70-80)               0.001670
gender|Male               0.001650
A1Cresult|None           0.001636
gender|Female              0.001619
class|>30                 0.001573
discharge_disposition_id|1 0.001569
race|Caucasian            0.001503
payer_code|MC              0.001457
discharge_disposition_id|18 0.001454
age|[80-90)               0.001407
age|[60-70)               0.001377
glyburide|Up                0.001376
admission_type_id|1         0.001370
diag|428                  0.001364
dtype: float64

correlated_features = correlation_matrix[abs(correlation_matrix) >
0.05].index.tolist()

important_features = importances_sorted[importances_sorted >
0.005].index.tolist()

selected_features = list(set(correlated_features) |
set(important_features))

if 'diabetesMed|Yes' in selected_features:
    selected_features.remove('diabetesMed|Yes')

print("Features with absolute correlation > 0.05:")
print(correlated_features)
print(f"\nNumber of selected features: {len(correlated_features)}")
print("\nFeatures with importance > 0.005:")
print(important_features)
print(f"\nNumber of selected features: {len(important_features)}")
print("\nSelected features (meeting both criteria):")

```

```

print(selected_features)
print(f"\nNumber of selected features: {len(selected_features)}")

Features with absolute correlation > 0.05:
['time_in_hospital', 'num_medications', 'diag|250.02', 'change|No',
'diabetesMed|Yes', 'discharge_disposition_id|18', 'payer_code|Missing',
'A1Cresult|>8', 'A1Cresult|None', 'metformin|No',
'metformin|Steady', 'metformin|Up', 'repaglinide|No', 'repaglinide|Steady',
'glimepiride|No', 'glimepiride|Steady', 'glipizide|No',
'glipizide|Steady', 'glyburide|No', 'glyburide|Steady', 'pioglitazone|No',
'pioglitazone|Steady', 'rosiglitazone|No', 'rosiglitazone|Steady',
'insulin|Down', 'insulin|No', 'insulin|Steady', 'insulin|Up',
'class|NO']

Number of selected features: 29

Features with importance > 0.005:
['insulin|No', 'insulin|Steady', 'change|No', 'metformin|No',
'metformin|Steady', 'glipizide|No', 'glipizide|Steady', 'glyburide|No',
'glyburide|Steady', 'insulin|Down', 'insulin|Up', 'pioglitazone|No',
'pioglitazone|Steady', 'num_medications', 'glimepiride|Steady',
'glimepiride|No', 'rosiglitazone|No', 'rosiglitazone|Steady',
'num_lab_procedures', 'time_in_hospital']

Number of selected features: 20

Selected features (meeting both criteria):
['metformin|Steady', 'repaglinide|Steady', 'time_in_hospital',
'num_medications', 'A1Cresult|None', 'A1Cresult|>8', 'diag|250.02',
'glipizide|Steady', 'glimepiride|No', 'metformin|Up', 'rosiglitazone|No',
'insulin|Steady', 'glimepiride|Steady', 'num_lab_procedures',
'payer_code|Missing', 'rosiglitazone|Steady', 'pioglitazone|No',
'glyburide|No', 'pioglitazone|Steady', 'insulin|Up', 'metformin|No',
'discharge_disposition_id|18', 'class|NO', 'glipizide|No', 'insulin|No',
'insulin|Down', 'glyburide|Steady', 'repaglinide|No', 'change|No']

Number of selected features: 29

```

We have chosen to select features that have either correlation higher than 0.05 or importance higher than 0.005. This gives us 29 features

```

X = df[selected_features]
y = df['diabetesMed|Yes']
dataset = pd.concat([X, y], axis=1)
dataset.head()

  metformin|Steady  repaglinide|Steady  time_in_hospital
num_medications \
0                  0                   0            0.272727

```

```

0.470588
1          0          0          0.181818
0.264706
2          1          0          0.090909
0.205882
3          0          0          0.000000
0.323529
4          1          0          0.181818
0.647059

      A1Cresult|None  A1Cresult|>8  diag|250.02  glipizide|Steady \
0          1          0          0          1
1          1          0          0          0
2          1          0          0          0
3          1          0          0          0
4          1          0          0          0

      glimepiride|No  metformin|Up   ...  metformin|No \
0          1          0          ...          1
1          1          0          ...          1
2          1          0          ...          0
3          1          0          ...          1
4          1          0          ...          0

      discharge_disposition_id|18  class|NO  glipizide|No  insulin|No \
0          0          1          0          0
1          0          1          1          0
2          0          1          1          0
3          0          1          1          1
4          0          1          1          1

      insulin|Down  glyburide|Steady  repaglinide|No  change|No
diabetesMed|Yes
0          0          0          1          0
1
1          0          1          1          0
2          0          0          1          0
1
3          0          0          1          1
0
4          0          1          1          0
1

[5 rows x 30 columns]

class_counts = dataset['diabetesMed|Yes'].value_counts()
print("Class distribution:\n", class_counts)

# --- Plot distribution ---


```

```

plt.figure(figsize=(6,4))
class_counts.plot(kind='bar', color=['skyblue','salmon'])
plt.title('Distribution of diabetesMed|Yes')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.show()

Class distribution:
diabetesMed|Yes
1    78363
0    23403
Name: count, dtype: int64

```



Dataset is imbalanced, therefore we will punish misclassification of the minority group more heavily.

```

dataset.to_csv('../processed_dataset/processed_dataset.csv',
index=False)

```

Experiments

After testing the model with this dataframe, the results were too accurate, meaning there were likely features that correlated too highly with the predicted feature, meaning the model could (and did) use these features to quickly and accurately get the correct prediction.

After further research we found out that these features were all types of diabetes medication: insulin, metformin, glyburide, glipizide, glyburide.metformin, nateglinide, pioglitazone, rosiglitazone, glimepiride, repaglinide, acarbose, miglitol, troglitazone, tolazamide, tolbutamide, chlorpropamide, sitagliptin.

Perform another data selection with different features. (The same as in dataloader.py)

```
import pandas as pd
import numpy as np
import torch
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
import openml

dataset_id = 45069
test_size = 0.2
random_state = 42
```

Load the data.

```
dataset = openml.datasets.get_dataset(dataset_id)
X, y, _, _ = dataset.get_data(
    dataset_format="dataframe",
    target=dataset.default_target_attribute
)

if y is not None:
    df = pd.concat([X, y], axis=1)
else:
    df = X

print(f"Loaded {df.shape[0]} records with {df.shape[1]} features")
Loaded 101766 records with 48 features
```

Drop unnecessary columns.

```
columns_to_drop = ['examide', 'citoglipton', 'weight']

df = df.drop(columns=columns_to_drop)
print(f"Dropped columns: {columns_to_drop}")

Dropped columns: ['examide', 'citoglipton', 'weight']
```

Convert to categorical columns.

```
columns_to_convert = [
    'admission_type_id',
    'discharge_disposition_id',
    'admission_source_id',
    'class',
    'diabetesMed',
]
for col in columns_to_convert:
    if col in df.columns:
        df[col] = df[col].astype('category')
```

Handle missing (NaN) values by replacing them with a new category called 'Missing'.

```
categorical_cols_with_missing = ['payer_code', 'medical_specialty']

for col in categorical_cols_with_missing:
    if col in df.columns:
        if pd.api.types.is_categorical_dtype(df[col].dtype):
            if 'Missing' not in df[col].cat.categories:
                df[col] = df[col].cat.add_categories('Missing')
            df[col] = df[col].fillna('Missing')

C:\Users\maxik\AppData\Local\Temp\ipykernel_38132\1295604742.py:5:
DeprecationWarning: is_categorical_dtype is deprecated and will be
removed in a future version. Use isinstance(dtype,
pd.CategoricalDtype) instead
    if pd.api.types.is_categorical_dtype(df[col].dtype):
```

Create unified diagnosis columns using all 3 current diagnosis columns.

```
all_diag_codes = pd.concat([
    df['diag_1'],
    df['diag_2'],
    df['diag_3']
]).unique()
all_diag_codes = [code for code in all_diag_codes if pd.notna(code)]

print(f"Creating {len(all_diag_codes)} diagnosis binary columns...")

diag_cols_dict = {
    f'diag|{code}': (
        (df['diag_1'] == code) |
        (df['diag_2'] == code) |
        (df['diag_3'] == code)
    ).astype(int)
    for code in all_diag_codes
}
```

```
diag_cols_df = pd.DataFrame(diag_cols_dict, index=df.index)
df = df.drop(columns=['diag_1', 'diag_2', 'diag_3'])
df = pd.concat([df, diag_cols_df], axis=1)
```

Creating 915 diagnosis binary columns...

Cap outliers using the IQR method.

```
columns_to_cap = [
    'number_diagnoses',
    'num_lab_procedures',
    'num_medications',
    'num_procedures',
    'time_in_hospital'
]

for col in columns_to_cap:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    df[col] = df[col].apply(
        lambda x: lower_bound if x < lower_bound else (
            upper_bound if x > upper_bound else x
        )
    )

print(f"Capped outliers for {len(columns_to_cap)} columns")
```

Capped outliers for 5 columns

One-hot encode the columns.

```
categorical_cols =
df.select_dtypes(include=['category']).columns.tolist()

cols_to_drop_first = [
    col for col in categorical_cols
    if df[col].nunique() == 2
]

cols_to_keep_all = [
    col for col in categorical_cols
    if df[col].nunique() > 2
]
```

```

df = pd.get_dummies(
    df,
    columns=cols_to_drop_first,
    drop_first=True,
    prefix_sep='|'
)
df = pd.get_dummies(
    df,
    columns=cols_to_keep_all,
    drop_first=False,
    prefix_sep='|'
)

encoded_cols = [
    col for col in df.columns
    if df[col].dtype != 'float'
]
df[encoded_cols] = df[encoded_cols].astype(int)

print(f"One-hot encoded: {df.shape[1]} columns")
One-hot encoded: 1158 columns

```

Select features based on a correlation_threshold = 0.05 and an importance_threshold = 0.005

```

correlation_threshold = 0.05
importance_threshold = 0.005

y = df['diabetesMed|Yes']
diabetes_cols = [c for c in df.columns if
c.startswith('diabetesMed|')]
X = df.drop(columns=diabetes_cols)

correlation_matrix = df.corrwith(df['diabetesMed|Yes'])
correlated_features = correlation_matrix[
    abs(correlation_matrix) > correlation_threshold
].index.tolist()

X_train_fs, X_test_fs, y_train_fs, y_test_fs = train_test_split(
    X, y, test_size=test_size, random_state=random_state
)

rf_model = RandomForestClassifier(
    n_estimators=200,
    random_state=random_state,
    n_jobs=-1
)
rf_model.fit(X_train_fs, y_train_fs)

importances = pd.Series(rf_model.feature_importances_,

```

```

index=X.columns)
important_features = importances[
    importances > importance_threshold
].index.tolist()

selected_features = list(set(correlated_features) |
set(important_features))
selected_features = [f for f in selected_features if not
f.startswith('diabetesMed|')]

print(f"Selected {len(selected_features)} features")
print(f" - Correlated features: {len(correlated_features)}")
print(f" - Important features: {len(important_features)}")

Selected 29 features
- Correlated features: 29
- Important features: 20

```

Prepare data and drop duplicates.

```

X = df[selected_features]
y = df['diabetesMed|Yes']

print(f"Shape before dropping duplicates: {X.shape}")
X = X.drop_duplicates()
y = y.loc[X.index]
print(f"Shape after dropping duplicates: {X.shape}")

Shape before dropping duplicates: (101766, 29)
Shape after dropping duplicates: (92766, 29)

```

Create train-test split.

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=test_size,
    random_state=random_state,
    stratify=y
)

X_train = torch.tensor(X_train.to_numpy(), dtype=torch.float32)
y_train = torch.tensor(y_train.to_numpy(),
dtype=torch.float32).unsqueeze(1)
X_test = torch.tensor(X_test.to_numpy(), dtype=torch.float32)
y_test = torch.tensor(y_test.to_numpy(),
dtype=torch.float32).unsqueeze(1)

print("\nFinal tensor shapes:")
print(f" X_train: {X_train.shape}")

```

```

print(f" X_test: {X_test.shape}")
print(f" y_train: {y_train.shape}")
print(f" y_test: {y_test.shape}")

Final tensor shapes:
X_train: torch.Size([74212, 29])
X_test: torch.Size([18554, 29])
y_train: torch.Size([74212, 1])
y_test: torch.Size([18554, 1])

```

Select features based on correlation and RandomForest importance using only training data.

```

def select_features_train(X_train_df: pd.DataFrame, y_train_s: pd.Series,
                         correlation_threshold: float = 0.05,
                         importance_threshold: float = 0.005) -> list:

    std = X_train_df.std(numeric_only=True)
    non_constant_cols = std[std > 0].index
    dropped_zero_var = X_train_df.shape[1] - len(non_constant_cols)
    if dropped_zero_var > 0:
        print(f"Dropped {dropped_zero_var} zero-variance columns from train before correlation")
    X_train_nc = X_train_df[non_constant_cols]

    with np.errstate(invalid='ignore', divide='ignore'):
        corr =
X_train_nc.corrwith(y_train_s).abs().fillna(0.0).sort_values(ascending=False)
        corr_features = corr[corr > correlation_threshold].index.tolist()

        rf_model = RandomForestClassifier(n_estimators=200,
random_state=random_state, n_jobs=-1)
        rf_model.fit(X_train_nc, y_train_s)
        importances = pd.Series(rf_model.feature_importances_,
index=X_train_nc.columns)
        imp_features = importances[importances > importance_threshold].index.tolist()

        selected = list(set(corr_features) | set(imp_features))
        print(f"Selected {len(selected)} features (train-only)")
        print(f" - Correlated: {len(corr_features)} Important: {len(imp_features)}")
    return selected

```

Load and preprocess the data.

```

# Define target and remove leaky columns
if 'diabetesMed|Yes' not in df.columns:
    raise ValueError("Target column 'diabetesMed|Yes' not found after
encoding.")

y_full = df['diabetesMed|Yes']

# Drop any one-hot columns derived from specific diabetes medications
# or the target/change
leaky_prefixes = [
    'diabetesMed', 'change', 'insulin', 'metformin', 'glyburide',
    'glipizide',
    'glyburide.metformin', 'nateglinide', 'pioglitazone',
    'rosiglitazone',
    'glimepiride', 'repaglinide', 'acarbose', 'miglitol',
    'troglitazone',
    'tolazamide', 'tolbutamide', 'chlorpropamide', 'sitagliptin'
]
leaky_cols = [c for c in df.columns if any(c.startswith(prefix + '| '))
for prefix in leaky_prefixes]
if leaky_cols:
    print(f"Dropping medication-related columns: {len(leaky_cols)}")
X_full = df.drop(columns=leaky_cols + ['diabetesMed|Yes'])

# Train/Test split
X_train_df, X_test_df, y_train_s, y_test_s = train_test_split(
    X_full, y_full, test_size=test_size, random_state=random_state,
    stratify=y_full
)

# Feature selection using only training data
selected = select_features_train(X_train_df, y_train_s)
selected_features = selected

# Apply selected features to both splits
X_train_df = X_train_df[selected]
X_test_df = X_test_df[selected]

# Optionally drop duplicates on train only
print(f"Shape before dropping duplicates (train): {X_train_df.shape}")
X_train_df = X_train_df.drop_duplicates()
y_train_s = y_train_s.loc[X_train_df.index]
print(f"Shape after dropping duplicates (train): {X_train_df.shape}")

# Convert to tensors
X_train = torch.tensor(X_train_df.to_numpy(), dtype=torch.float32)
y_train = torch.tensor(y_train_s.to_numpy(),
                      dtype=torch.float32).unsqueeze(1)
X_test = torch.tensor(X_test_df.to_numpy(), dtype=torch.float32)
y_test = torch.tensor(y_test_s.to_numpy()),

```

```

        dtype=torch.float32).unsqueeze(1)

n_features = X_train.shape[1]

print("*"*80)
print(f"PREPROCESSING COMPLETE - {n_features} features selected
(train-only selection, no leakage)")
print("*"*80)

Dropping leaky medication-related columns: 59
Dropped 19 zero-variance columns from train before correlation
Selected 47 features (train-only)
- Correlated: 8 Important: 45
Shape before dropping duplicates (train): (81412, 47)
Shape after dropping duplicates (train): (81406, 47)
=====
=====
PREPROCESSING COMPLETE - 47 features selected (train-only selection,
no leakage)
=====
=====
```

Final dataframe inspection.

```

print(f"\n{'='*80}")
print("DATA SHAPES")
print(f"{'='*80}")
print(f"X_train: {X_train.shape}")
print(f"X_test: {X_test.shape}")
print(f"y_train: {y_train.shape}")
print(f"y_test: {y_test.shape}")
print(f"Number of features (n_input): {n_features}")

# Get feature names
if selected_features is not None:
    feature_names = selected_features
    print(f"\n{'='*80}")
    print(f"SELECTED FEATURES ({len(feature_names)} total)")
    print(f"{'='*80}")
    for i, feat in enumerate(feature_names, 1):
        print(f"{i:2d}. {feat}")
else:
    print("\nWarning: Could not retrieve feature names")

# Show class distribution
print(f"\n{'='*80}")
print("TARGET DISTRIBUTION")
print(f"{'='*80}")
import torch
```

```

if isinstance(y_train, torch.Tensor):
    y_train_np = y_train.cpu().numpy()
    y_test_np = y_test.cpu().numpy()
else:
    y_train_np = y_train
    y_test_np = y_test

train_pos = int(y_train_np.sum())
train_total = len(y_train_np)
test_pos = int(y_test_np.sum())
test_total = len(y_test_np)

print(f"Train: {train_pos}/{train_total} ({100*train_pos/train_total:.1f}%) positive")
print(f"Test: {test_pos}/{test_total} ({100*test_pos/test_total:.1f}%) positive")

# Show basic stats for X_train
print(f"\n{'='*80}")
print("FEATURE STATISTICS (Training Set)")
print(f"{'='*80}")

if isinstance(X_train, torch.Tensor):
    X_train_np = X_train.cpu().numpy()
else:
    X_train_np = X_train

df_stats = pd.DataFrame({
    'Feature': feature_names if 'feature_names' in locals() else
    [f"Feature_{i}" for i in range(n_features)],
    'Mean': X_train_np.mean(axis=0),
    'Std': X_train_np.std(axis=0),
    'Min': X_train_np.min(axis=0),
    'Max': X_train_np.max(axis=0),
    'Non-zero %': 100 * (X_train_np != 0).sum(axis=0) /
    len(X_train_np),
})
print(f"\n{'='*80}")
print("Top 20 features:")
print(df_stats.head(20).to_string(index=False))

print(f"\n{'='*80}")
print("INSPECTION COMPLETE")
print(f"{'='*80}")

```

=====
=====
DATA SHAPES

```
=====
=====
X_train: torch.Size([81406, 47])
X_test: torch.Size([20354, 47])
y_train: torch.Size([81406, 1])
y_test: torch.Size([20354, 1])
Number of features (n_input): 47

=====
=====
SELECTED FEATURES (47 total)
=====
=====

1. time_in_hospital
2. number_emergency
3. age|[70-80)
4. discharge_disposition_id|18
5. diag|414
6. discharge_disposition_id|3
7. payer_code|Missing
8. admission_type_id|3
9. class|>30
10. diag|403
11. age|[40-50)
12. num_procedures
13. diag|276
14. age|[80-90)
15. diag|250.02
16. age|[60-70)
17. race|AfricanAmerican
18. diag|428
19. gender|Male
20. race|Caucasian
21. age|[50-60)
22. diag|427
23. diag|780
24. gender|Female
25. num_medications
26. admission_type_id|2
27. discharge_disposition_id|1
28. diag|786
29. diag|496
30. diag|250
31. A1Cresult|>8
32. discharge_disposition_id|6
33. number_inpatient
34. medical_specialty|Missing
35. medical_specialty|Family/GeneralPractice
36. medical_specialty|InternalMedicine
```

```
37. A1Cresult|None  
38. number_outpatient  
39. class|N0  
40. number_diagnoses  
41. diag|401  
42. admission_source_id|7  
43. admission_source_id|1  
44. diag|599  
45. payer_code|MC  
46. num_lab_procedures  
47. admission_type_id|1
```

TARGET DISTRIBUTION

```
Train: 62685/81406 (77.0%) positive  
Test: 15673/20354 (77.0%) positive
```

FEATURE STATISTICS (Training Set)

Top 20 features:

	Feature	Mean	Std	Min	Max	Non-zero	%
	time_in_hospital	4.360784	2.887990	1.0	12.0	100.000000	
	number_emergency	0.196226	0.924113	0.0	76.0	11.150284	
	age [70-80)	0.256492	0.436697	0.0	1.0	25.649215	
discharge_disposition_id 18	0.036201	0.186791	0.0	1.0	3.620126		
	diag 414	0.120937	0.326054	0.0	1.0	12.093703	
discharge_disposition_id 3	0.137373	0.344241	0.0	1.0	13.737317		
	payer_code Missing	0.394909	0.488831	0.0	1.0	39.490947	
admission_type_id 3	0.185773	0.388923	0.0	1.0	18.577255		
	class >30	0.348328	0.476441	0.0	1.0	34.832813	
diag 403	0.055868	0.229667	0.0	1.0	5.586812		
	age [40-50)	0.095521	0.293933	0.0	1.0	9.552121	
num_procedures	1.285028	1.577680	0.0	5.0	54.090608		
	diag 276	0.127312	0.333323	0.0	1.0	12.731248	
age [80-90)	0.168956	0.374713	0.0	1.0	16.895561		
	diag 250.02	0.040513	0.197159	0.0	1.0	4.051298	
age [60-70)	0.220303	0.414451	0.0	1.0	22.030317		
	race AfricanAmerican	0.188414	0.391042	0.0	1.0	18.841363	
diag 428	0.171412	0.376869	0.0	1.0	17.141243		
	gender Male	0.463185	0.498643	0.0	1.0	46.318453	
race Caucasian	0.747832	0.434257	0.0	1.0	74.783186		

```
=====
INSPECTION COMPLETE
=====
```

The transformations showcased above are called before training the model. We load, transform data and select features using dataloader class and then we train the model with trainer and model classes.

Configuration

Model is configured using configuration file. These configuration files are contained inside .scratch/experiment_configs folder. For each experiment we create a file and then we run tools.main. This calls run command which logs the run to wandb.

Example of config file:

```
n_input: 47 layers: [512, 256, 128, 32, 8, 1]
activations: [ReLU, ReLU, ReLU, ReLU, ReLU, null]
output_activation: null
dropout: 0.1
batch_norm: true
skip_connections: false
bottleneck: true

training:
  epochs: 100
  lr: 0.0005
  batch_size: 256
  verbose: true
  patience: 10
  min_delta: 0.001
```

Explanation:

n_input - Number of input features fed into the first layer of the network.
layers - List defining the number of neurons in each hidden and output layer.
activations - Activation functions applied after each layer (e.g., ReLU, GELU).
output_activation - Final activation used for the output layer (e.g., Sigmoid for binary tasks).
dropout - Fraction of neurons randomly turned off during training to prevent overfitting.
batch_norm - Whether to apply Batch Normalization for stabilizing training.
skip_connections - Lets the model reuse information from earlier layers to learn better and faster.
bottleneck - Adds smaller intermediate layers to force the network to learn compact representations.

epochs - Maximum number of full passes through the training dataset.
lr - Learning rate controlling how much weights are updated per step.
batch_size - Number of samples processed together before an optimizer update.
verbose - Whether to print progress and loss values during training.

patience - Number of epochs to wait for improvement before early stopping.
min_delta - Minimum change in validation loss to be considered an improvement.

Experiment tracking

For our experiment tracking we use wandb. We call run from main and this saves the results of the tests.

Picture of runs

Confusion table

Learning process1

Learning process2

We can also compare the last 10 runs to see if we are improving.

Learning processes comparison

Experiment

For our experiments we have tried to change batch sizes, number of layers, applying optimization techniques - Dropout, Skip Connections, Bottleneck Layers. We have also tried changing the learning rate, patience and min_delta.

With prepared data we were able to train the model in 12 epochs to:

```
test/accuracy:0.9956273951066128  
test/f1:0.9971526378091308  
test/precision:1  
test/roc_auc:0.997850680470756  
test/sensitivity:0.9943214445224272  
test/specificity:1
```

Confusion Matrix

Confusion Matrix

We have used this configuration file:

```
n_input: 29 layers: [128, 64, 32, 1]  
activations: [ReLU, ReLU, ReLU, null]  
output_activation: null  
dropout: 0.3  
batch_norm: true  
skip_connections: false  
bottleneck: true  
  
training:  
    epochs: 100  
    lr: 0.001  
    batch_size: 256
```

```
verbose: true  
patience: 10  
min_delta: 0.001
```

Process of training:

Learning processes comparison

This result was expected, the high precision was achieved by including data about medications, which closely relates to if patient is using medications. If person is using any concrete medication for diabetes they are using medication for diabetes for sure.

To have a model with relevant purpose we have removed these medications from our data. We have removed 'diabetesMed', 'change', 'insulin', 'metformin', 'glyburide', 'glipizide', 'glyburide.metformin', 'nateglinide', 'pioglitazone', 'rosiglitazone', 'glimepiride', 'repaglinide', 'acarbose', 'miglitol', 'troglitazone', 'tolazamide', 'tolbutamide', 'chlorpropamide', 'sitagliptin' which are medications that treat diabetes (change refers to change in medications). Next we did experiments using these new data.

For these we have found 3 models that can be thought of as the best models, first is best overall, second one is best when we want to minimize detection of diabetes with non-diabetic people and third when we want to prioritize detection of diabetes with diabetic people.

1st model

For this model were able to achieve in 100 epochs:

```
test/accuracy:0.6550555173430284  
test/f1:0.7451265110538353  
test/precision:0.8643254168772108  
test/roc_auc:0.7076594152879849  
test/sensitivity:0.654820391756524  
test/specificity:0.6558427686391797  
train/loss:0.27352656585709106
```

This model has accuracy of 65% and loss of 0.2735. We were not able to find better model.

Confusion matrix:

Confusion Matrix

We have used this configuration file:

```
n_input: 47 layers: [90, 29, 1]  
activations: [ReLU, ReLU, null]  
output_activation: null  
dropout: 0.2  
batch_norm: true  
skip_connections: false  
bottleneck: true
```

training:

epochs: 100

```
lr: 0.01  
batch_size: 128  
verbose: true  
patience: 10  
min_delta: 0.001
```

Process of training:

Learning processes comparison

2nd model

For this model were able to achieve in 54 epochs:

```
test/accuracy:0.47150437260489336  
test/f1:0.5008121026497749  
test/precision:0.9183117767188564  
test/roc_auc:0.7211868025425039  
test/sensitivity:0.3442863523256556  
test/specificity:0.8974578081606495  
train/loss:0.26762372878549984
```

This model has low accuracy but high precision.

Confusion matrix:

Confusion Matrix

We have used this configuration file:

```
n_input: 80 layers: [128,64,32,16,1]  
activations: [ReLU, ReLU, ReLU, ReLU,null]  
output_activation: null  
dropout: 0.1  
batch_norm: true  
skip_connections: false  
bottleneck: true
```

training:

```
epochs: 100  
lr: 0.0001  
batch_size: 128  
verbose: true  
patience: 30  
min_delta: 0.001
```

Process of training:

Learning processes comparison

3rd model

For this model were able to achieve in 100 epochs:

```
test/accuracy:0.7711997641741181  
test/f1:0.8607231510003888  
test/precision:0.8100653006079712  
test/roc_auc:0.7224563738997474  
test/sensitivity:0.9181394755311684  
test/specificity:0.2792138431958983  
train/loss:0.2656220795832715
```

This model has high sensitivity.

Confusion matrix:

Confusion Matrix

We have used this configuration file:

```
n_input: 47 layers: [90, 29, 1]  
activations: [ReLU, ReLU, null]  
output_activation: null  
dropout: 0.2  
batch_norm: true  
skip_connections: false  
bottleneck: true
```

training:

```
epochs: 100  
lr: 0.01  
batch_size: 128  
verbose: true  
patience: 10  
min_delta: 0.001
```

Process of training:

Learning processes comparison

